

Lecture 1: Introduction, Polynomial Curve

H1 Fitting, and Probability Theory - 13/01/20

H2 Notations



using *Classifying Hand Written Digits* as example

- A training set of N digits
- Each digit, i , is an image, representing as an **input vector of pixel values** x_i
- The category of each digit, i , is known and expresses as **target vector** t_i
- ML algorithm outputs function $y(x)$, which can take new digit input x and output vector y , which is a **guess** of the target t . The precise form of $y(x)$ is determined during the training phases.
- The ability to categorise new examples that differ from those used for training is called **generalisation**

H2 Supervised Learning

Problems are ones where the data contains both input and corresponding target vectors.

- **Classification**
- **Regression**

The inputs may be **pre-processed** to reduce variability in the inputs.

H2 Unsupervised Learning


Problems are ones where the data contains only input vectors but no targets.

- **Clustering** - discovering groups of similar examples
- **Density Estimation** - learning how data is distributed
- **Dimensionality Reduction** - representing high dimensional data with just a few variables

H2 Reinforcement Learning

Problems are ones that interact with an environment by choosing actions and observing changes in state. Actions must act to maximise a **reward signal**. Optimal actions are discovered by **trials and errors**.

H2 Polynomial Curve Fitting



- **Training inputs** $\mathbf{x} = (x_1, \dots, x_N)^T$
- **Training targets** $\mathbf{t} = (t_1, \dots, t_N)^T$

This is **synthetic data** - we know how it originated

- Each x_i is sampled uniformly from $[0, 1]$
- Each $t_i = \sin(2\pi x_i) + (\text{Gaussian Noise})$

Data tends to have an underlying regularity or structure obscured by noise. Noise can be:

- **intrinsically stochastic** (random)
- resulted of **unobserved** sources of **variability**

H3 Aim

- Predict a target \hat{t} for an unseen input \hat{x}
- Discover the **underlying structure**
- Sparate it from the **noise**

H3 Fitting with Linear Model

$$y(x, \mathbf{w}) = w_0 + w_1 x + w_2 x^2 + \dots + w_M x^M = \sum_{j=0}^M w_j x^j$$

- M is the **order of the polynomial**
- **Polynomial coefficients** w_0, \dots, w_M are collected into vector \mathbf{w}
- $y(x, \mathbf{w})$ is non-linear in x , but it is linear in \mathbf{w} and so we call this a **linear model**

We estimate values for \mathbf{w} by fitting the function to training data. Fit the function by **minising** an **error function**

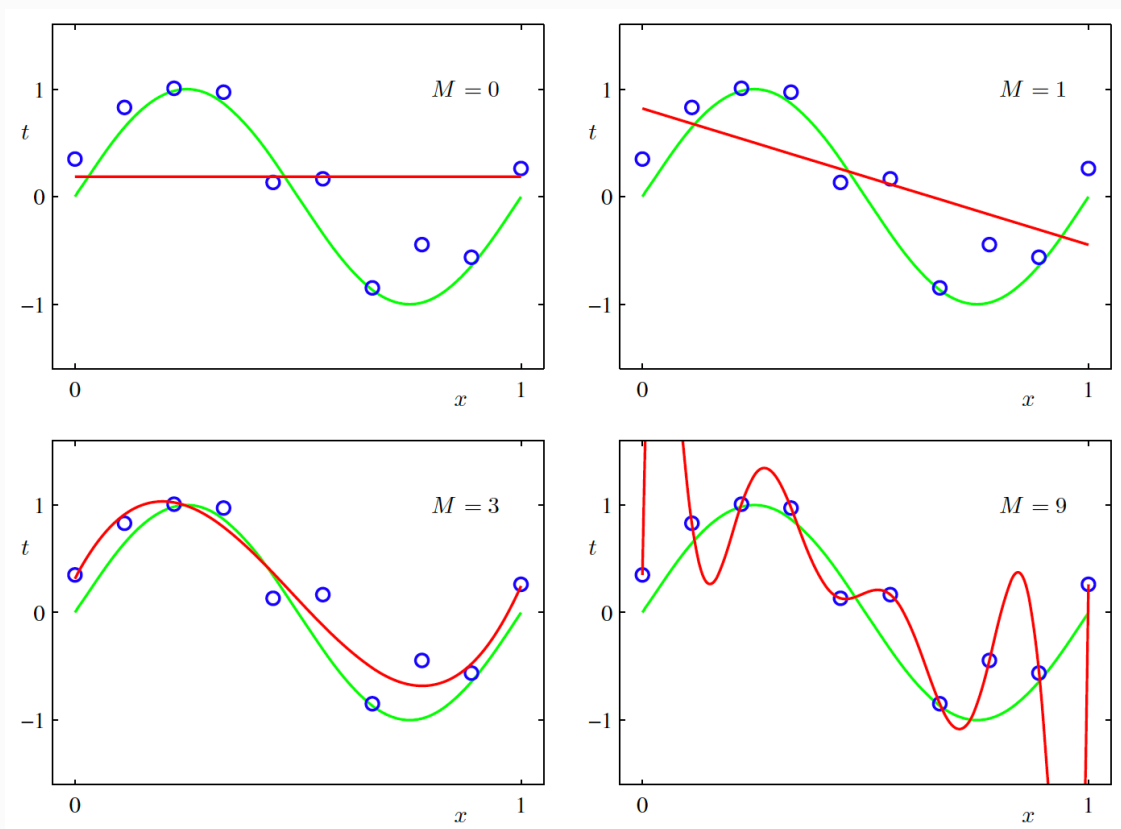
H3 Error Function

A widely used error function is the **Sum of Square Errors**

$$E(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N [y(x_n, \mathbf{w}) - t_n]^2$$

- **Best Fit** $\mathbf{w}^* = \arg \min_{\mathbf{w}} E(\mathbf{w})$
- **Perfect Fit** if $E(\mathbf{w}^*) = 0$
- Bigger differences are increasingly **penalised**

H3 Finding the Best Polynomial Degree



Choosing the best M is an example of **Model Selection**

- Small values of M give a poor fit
- Large values of M appear to **over-fit** - *capture the noise* rather than underlying structure

H2 Evaluating Fit and Regularisation

We need an objective way to test our fit

Root Mean Squared Error (RMSE)


$$E_{RMS} = \sqrt{\frac{2}{N} E(\mathbf{w}^*)}$$

Comparable for different amounts of data

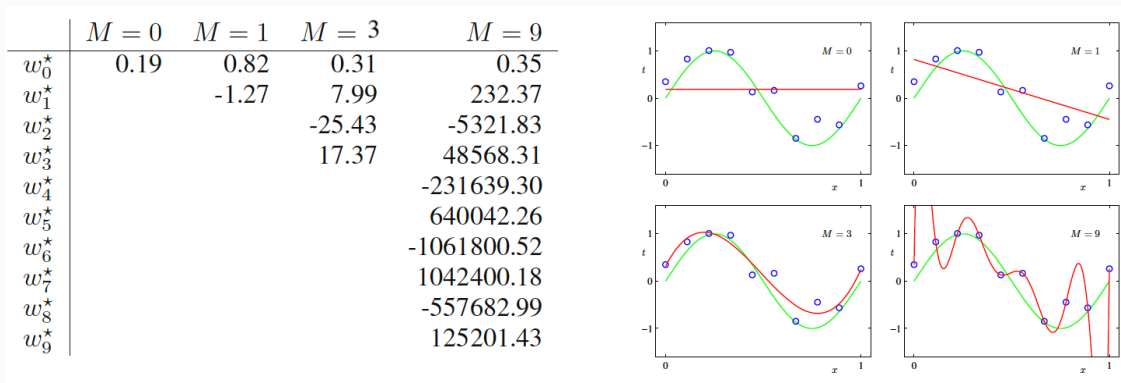
H2 Avoiding Over-fitting

H3 Indirect Evidence of Over-fitting

- Dramatic increase of E_{RMS} of training set and the difference between the E_{RMS} of training set and testing set as degree gets larger
- Magnitude of \mathbf{w}_i^* is extremely large



Over-fitting means we fail to **generalise** to un seen data



For $M = 9$, the magnitude of some w_i^* are very large, and the model makes some extreme predictions

Dilemma: Complex Models (more expressive) v. Over-fitting

H3 **Solution 1: Use More Data**

H3 **Solution 2: Regularisation**

Using a new **error function** that **penalises** extreme parameter values

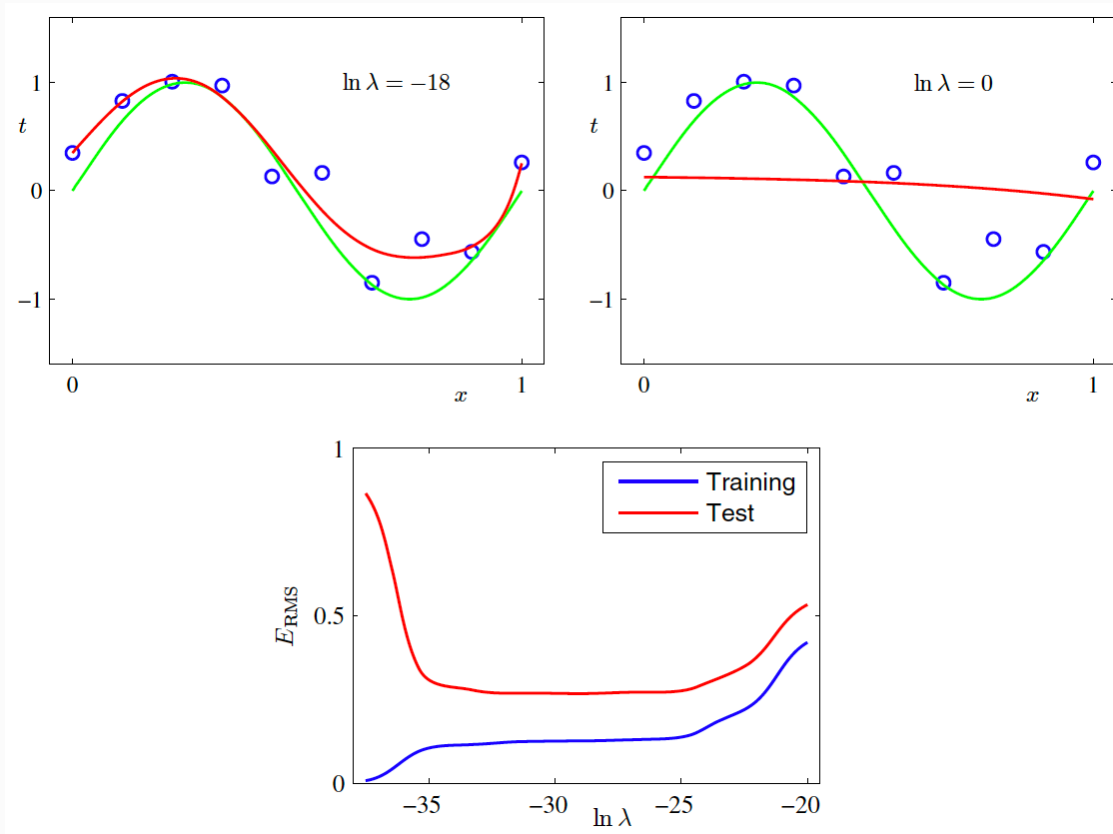
$$\tilde{E}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N [y(x_n, \mathbf{w}) - t_n]^2 + \frac{\lambda}{2}$$

Where

$$\|\mathbf{w}\|^2 = \mathbf{w}^T \mathbf{w} = w_0^2 + w_1^2 + \dots + w_M^2$$

Minimising **error function**

$$\mathbf{w}^* = \arg \min_w \tilde{E}(\mathbf{w})$$



Regularisation appears to control the effective complexity of the model, and hence the degree of overfitting.