


H1 Lecture 3: Regression in Linear Models - 27/01/20

H2 Linear Regression

A simple **linear model** for vector inputs $\mathbf{x} \in \mathbb{R}^{D-1}$:

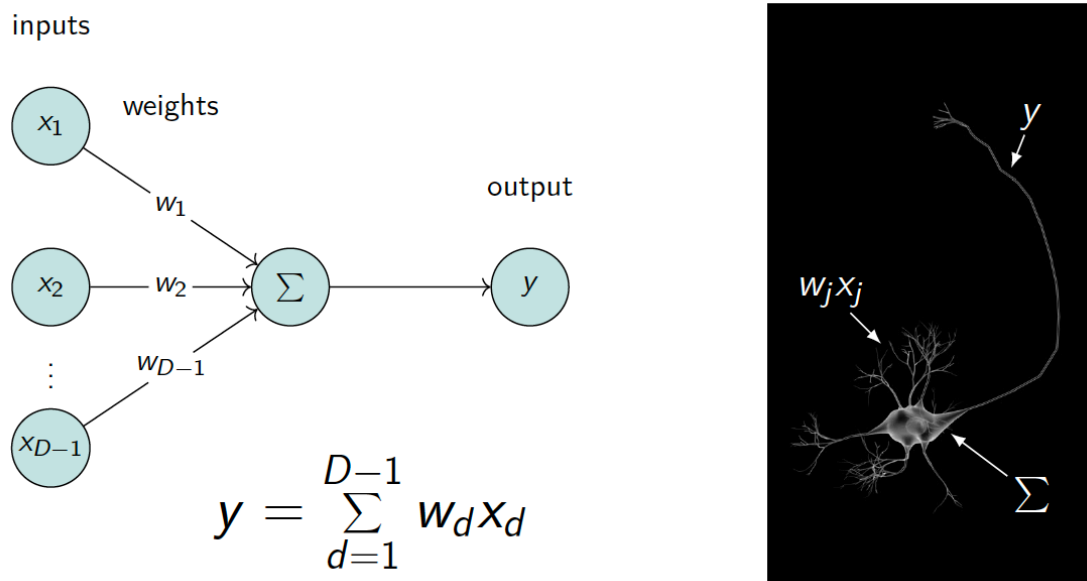
$$\begin{aligned} y(\mathbf{x}) &= \sum_{d=1}^{D-1} w_d x_d \\ &= (w_1 \quad \dots \quad w_{D-1}) \begin{pmatrix} x_1 \\ \vdots \\ x_{(D-1)} \end{pmatrix} \\ &= \begin{pmatrix} w_1 \\ \vdots \\ w_{(D-1)} \end{pmatrix}^T \begin{pmatrix} x_1 \\ \vdots \\ x_{(D-1)} \end{pmatrix} \\ &= \mathbf{w}^T \mathbf{x} \end{aligned}$$



$y = 0$ is defined as the **decision line**

H3 A Simple Neuron

We can think of this as equivalent to a simple neuron model called the perceptron:



H3 Bias Term

We introduce a **bias term**:

$$\begin{aligned} y(\mathbf{x}) &= w_0 + \sum_{d=1}^{D-1} w_d x_d \\ &= w_0 + \mathbf{w}^T \mathbf{x} \end{aligned}$$

For a point on the decision line:


$$\begin{aligned} y(\mathbf{x}) &= 0 \\ \frac{\mathbf{w}^T \mathbf{x}}{\|\mathbf{w}\|} &= -\frac{w_0}{\|\mathbf{w}\|} \end{aligned}$$

The bias can be absorbed into the vector:

$$\begin{aligned} y &= w_0 + \sum_{d=1}^{D-1} w_d x_d \\ &= \begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_{(D-1)} \end{pmatrix}^T \begin{pmatrix} 1 \\ x_1 \\ \vdots \\ x_{(D-1)} \end{pmatrix} \\ &= \mathbf{w}^T \mathbf{x} \end{aligned}$$

New definition of \mathbf{w} and \mathbf{x}

Model has D parameters $\{w_0, \dots, w_{(D-1)}\}$ (degrees of freedom).



H3 Solving the Linear Model

Each input is a vector $\mathbf{x}_n \in \mathbb{R}^D$, with corresponding target, $t \in \mathbb{R}$. We want to minimise the **sum-of-square errors**, with the **error function** being:

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \mathbf{w}^T \mathbf{x}_n)^2$$

Rewrite in matrix notation:

$$E_D(\mathbf{w}) = \frac{1}{2} (\mathbf{t} - \mathbf{X}\mathbf{w})^T (\mathbf{t} - \mathbf{X}\mathbf{w})$$

with $\mathbf{t} \in \mathbb{R}^N$ is our collected targets and $(N \times D)$ -matrix of inputs:

$$\mathbf{X} = \begin{pmatrix} x_{10} & x_{11} & \dots & x_{1(D-1)} \\ x_{20} & x_{21} & \dots & x_{2(D-1)} \\ \vdots & \vdots & \ddots & \vdots \\ x_{N0} & x_{N1} & \dots & x_{N(D-1)} \end{pmatrix}$$

Note that: each row vector \mathbf{x}_i^T is i th data input while each column vector is a set of data input $\tilde{\mathbf{x}}_j$ for j th demension

Minimise the error function $E_D(\mathbf{w})$ by differentiating and setting to zero:

$$\begin{aligned} \nabla_{\mathbf{w}} E_D(\mathbf{w}) &= \nabla_{\mathbf{w}} \left[\frac{1}{2} (\mathbf{t} - \mathbf{X}\mathbf{w})^T (\mathbf{t} - \mathbf{X}\mathbf{w}) \right] = 0 \\ &\quad -\mathbf{X}^T (\mathbf{t} - \mathbf{X}\mathbf{w}^*) = 0 \end{aligned}$$

Expanding brackets, rearraging, multipling by $(\mathbf{X}^T \mathbf{X})^{-1}$

$$\begin{aligned} \mathbf{X}^T \mathbf{X} \mathbf{w}^* &= \mathbf{X}^T \mathbf{t} \\ (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{X} \mathbf{w}^* &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t} \\ \mathbf{w}^* &= (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t} \end{aligned}$$


Most Likelihood Weights Linear Regression:

$$\mathbf{w}^* = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{t} = \mathbf{w}_{ML}$$

The Gradient Operator $\nabla_{\mathbf{w}}$

The gradient operator is vector of (partial) differential operations that gives direction of maximum ascent

$$\nabla_{\mathbf{w}} f = \frac{df}{d\mathbf{w}} = \left(\frac{\delta f}{\delta w_0}, \dots, \frac{\delta f}{\delta w_{(D-1)}} \right)^T$$



Properties:

- Gradient of dot product: $\nabla_{\mathbf{w}} \mathbf{w}^T \mathbf{v} = \mathbf{v}$
- Product rule: $\nabla_{\mathbf{x}} u(\mathbf{x})v(\mathbf{x}) = v \nabla_{\mathbf{x}} u + u \nabla_{\mathbf{x}} v$
- Chain rule: $\nabla_{\mathbf{z}} f(g(\mathbf{z})) = \frac{df}{dg} \nabla_{\mathbf{z}} g$

The Moore-Penrose Pseudo-Inverse

$$(\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T = \mathbf{A}^\dagger \in \mathbb{R}^{M \times N}$$

\mathbf{A}^\dagger is defined as the Moore-Penrose pseudo-inverse of matrix \mathbf{A} , which provides properties similar to the inverse of a square matrix for non-square matrix:

- Not a real inverse: $\mathbf{A} \mathbf{A}^\dagger \neq \mathbf{I}$
- Almost an inverse: $\mathbf{A} \mathbf{A}^\dagger \mathbf{A} = \mathbf{A}$
- If \mathbf{A} is square and invertible then $\mathbf{A}^\dagger = \mathbf{A}^{-1}$
- Can be problematic if $\mathbf{A}^T \mathbf{A}$ is (or close to) singular


H3 Geometric Intuition

$$\mathbf{X} = \begin{pmatrix} x_{10} & x_{11} & \dots & x_{1(D-1)} \\ x_{20} & x_{21} & \dots & x_{2(D-1)} \\ \vdots & \vdots & \ddots & \vdots \\ x_{N0} & x_{N1} & \dots & x_{N(D-1)} \end{pmatrix}$$

- Row Vector \mathbf{x}_i^T
- Column Vector $\tilde{\mathbf{x}}_j \in \mathbb{R}^N$
- \mathbf{t} is a vector in \mathbb{R}^N
- S is (sub-)space spanned by $\{\tilde{\mathbf{x}}_d\}$
- $\dim(S) \leq D$

Some of the data input might not be linearly independent

- $\mathbf{y} = \mathbf{X} \mathbf{w}^*$ is point in S closest to \mathbf{t}



H2 *k*NN for Regression

k-**Nearest Neighbours** (*k***NN**) assumes estimates $y(\mathbf{x})$ & $y(\mathbf{x}')$ are similar, when \mathbf{x} is close to \mathbf{x}' :

Predicts:

$$y(\mathbf{x}, k) = \frac{1}{k} \sum_{\mathbf{x}_i \in \mathbb{N}_k(\mathbf{x})} t_i$$

when $\mathbb{N}_k \mathbf{x}$ contains the k closest points to \mathbf{x}

In words, **predict for $y(\mathbf{x})$ the average target of the k nearest points**

- A closeness measure, e.g. Euclidean distance, is required
- Usually more common for **classification**

H3 Pseudocode

```

1: procedure kNN_REGRESSION( $\mathcal{D}$ ,  $\mathbf{x}$ ,  $k$ )
2:   #  $\mathcal{D} = \{(\mathbf{x}_n, t_n)\}_{n=1}^N$  is training data
3:   #  $\mathbf{x}$  is a test point,  $k$  is an integer
4:   sort  $\mathcal{D}$  by increasing distance  $d(\mathbf{x}_n, \mathbf{x})$ 
5:    $\mathbb{N}_k(\mathbf{x}) \leftarrow$  first  $k$  elements of sorted  $\mathcal{D}$ 
6:   return  $\frac{1}{k} \sum_{\mathbf{x}_n \in \mathbb{N}_k(\mathbf{x})} t_n$ 


```

H3 Insights


- No training phase
- Evaluation is expensive, sort is $O(N \log N)$
- Seems like it has one parameter, k

- Actually has $\frac{N}{k}$ effective parameters


H3 One-Dimensional k NN




$k = 1$




$k = 3$



$k = 5$




$k = 11$



Linear Regression

- Produces smooth function
- Stable fit
- Strong linear assumption restricts family of functions
- D parameters
- Low Variance, (potentially) High Bias



k NN Regression ($k = 3$)

- Weak assumptions
- Flexible functional form
- Unstable predictions (each estimate based on k obs.)
- $\frac{N}{k}$ effective parameters
- High Variance, Low Bias

H2 Linear Models

Consider a *simple linear regression* with vector inputs:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{d=1}^D w_d x_d$$

with vector input data, $\mathbf{x} = (x_1, \dots, x_D)$.

Linear in both **weights** and **input variables** x_i

Extending that to consider:

$$y(\mathbf{x}, \mathbf{w}) = w_o + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$$

where $\phi_j(\mathbf{x})$ are **basis functions**. For instance, a monomial function: $\phi_j(\mathbf{x}) = \sum_i x_i^j$

Also a **linear model** (linear in the weights, \mathbf{w})

Extended linear model:

$$y(\mathbf{x}, \mathbf{w}) = w_o + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x}) = \sum_{j=0}^{M-1} w_j \phi_j(\mathbf{x})$$

where $\phi_0(\mathbf{x}) = 1$

Rewrite in vector form as:

Linear Model Prediction

$$y(\mathbf{x}, \mathbf{w}) = \phi(\mathbf{x})^T \mathbf{w}$$

where $\phi(\mathbf{x})$ is our **feature vector**, defined as:


$$\phi(\mathbf{x}) = (\phi_0(\mathbf{x}), \phi_1(\mathbf{x}), \dots, \phi_{M-1}(\mathbf{x}))^T$$

H3 Example: Polynomial Basis Function

We can choose our basis functions very flexibly:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$$

- In 1d: $\phi_j(x) = x^j$
- Generally: $\phi_j(\mathbf{x}) = \prod_d x_d^{j_d}$
- **global functions** – change in one region of input space affects all others



H3 Example: Radial Basis Function

(Gaussian) Radial Basis Functions (RBF) are very common:

$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$$


- For 1d input x :

$$\phi_j(x) = \exp\left(-(x - \mu_j)^2 / 2s^2\right)$$

- Generally:

$$\phi_j(\mathbf{x}) = \exp\left[-\frac{(\mathbf{x} - \boldsymbol{\mu}_j)^T(\mathbf{x} - \boldsymbol{\mu}_j)}{2s^2}\right]$$

- RBF effects are **local**.



H3 Example: S-Shape Function

There are other choices of basis function for linear models:


$$y(\mathbf{x}, \mathbf{w}) = w_0 + \sum_{j=1}^{M-1} w_j \phi_j(\mathbf{x})$$

- The logistic function:

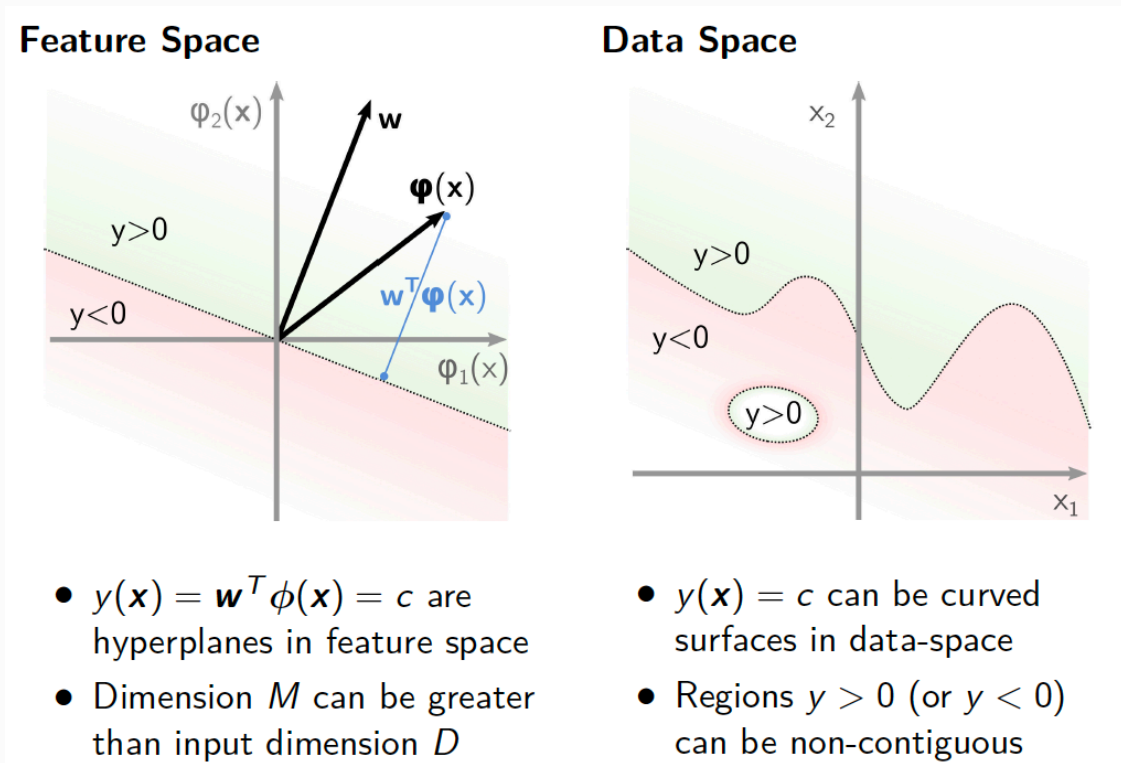
$$\phi_j(x) = \sigma\left(\frac{x - \mu_j}{s}\right)$$

$$\text{where } \sigma(a) = \frac{1}{1 + \exp(-a)}$$

- Similar functions, e.g. \tanh
- Multidimensional?



H3 Linear Models: Geometric Intuition



H2 Fitting Linear Models

Assuming target t given by deterministic component plus **Gaussian noise**:

$$t = y(\mathbf{x}; \mathbf{w}) + \epsilon$$

where $y(\mathbf{x}; \mathbf{w}) = \phi(\mathbf{x}_n)^T \mathbf{w}$ and $\epsilon \sim N(\cdot | 0, \beta^{-1})$

The **probability density** for target t :

$$p(t|\mathbf{x}, \mathbf{w}, \beta) = N(t|\phi(\mathbf{x}_n)^T \mathbf{w}, \beta^{-1})$$

The **conditional mean**:

$$E[t|\mathbf{x}, \mathbf{w}, \beta] = \int t p(t|\mathbf{x}, \mathbf{w}, \beta) dt = y(\mathbf{x}, \mathbf{w})$$

Collect all inputs together into data matrix $\mathbf{X} = (x_1^T, \dots, x_N^T)^T$ with vector of corresponding targets $\mathbf{t} = (t_1, \dots, t_N)^T$ now likelihood is :

$$p(\mathbf{t}|\mathbf{X}, \beta) = \prod_{n=1}^N N(t_n|\phi(\mathbf{x}_n)^T \mathbf{w}, \beta^{-1})$$


Taking log:

$$\begin{aligned} \ln p(\mathbf{t}|\mathbf{X}, \beta) &= \sum_{n=1}^N \ln N(t_n|\phi(\mathbf{x}_n)^T \mathbf{w}, \beta^{-1}) \\ &= \frac{N}{2} \ln \beta - \frac{N}{2} \ln(2\pi) - \beta E_D(\mathbf{w}) \end{aligned}$$


with **sum-of-squares error**:

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \phi(\mathbf{x}_n)^T \mathbf{w})^2$$


H3 Fitting with Different Basis Functions




Linear



Polynomial (degree 9)



RBF Basis Functions (9)



Logistic Basis Functions (9)

H3 Finding the Maximum Likelihood

Since the function is **quadratic** in \mathbf{w} , there is a single maximum. To maximise the likelihood, differentiate and set to zero:

$$\nabla_{\mathbf{w}} \ln p(\mathbf{t}|\mathbf{X}, \beta) = \beta \nabla_{\mathbf{w}} E_D(\mathbf{w}) = 0$$

Rewrite **error function** in matrix form, differentiate and set to zero:


$$\begin{aligned} E_D(\mathbf{w}) &= \frac{1}{2} (\mathbf{t} - \Phi \mathbf{w})^T (\mathbf{t} - \Phi \mathbf{w}) \\ \nabla_{\mathbf{w}} E_D(\mathbf{w}) &= -\Phi^T (\mathbf{t} - \Phi \mathbf{w}) = 0 \\ \Phi^T \Phi \mathbf{w} &= \Phi^T \mathbf{t} \\ \cancel{(\Phi^T \Phi)^{-1} \Phi^T \Phi} \mathbf{w} &= (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} = \Phi^\dagger \mathbf{t} \end{aligned}$$

where we have defined the design matrix as:

$$\Phi = \begin{pmatrix} \phi_0(\mathbf{x}_1) & \phi_1(\mathbf{x}_1) & \dots & \phi_{M-1}(\mathbf{x}_1) \\ \phi_0(\mathbf{x}_2) & \phi_1(\mathbf{x}_2) & \dots & \phi_{M-1}(\mathbf{x}_2) \\ \vdots & \vdots & \ddots & \vdots \\ \phi_0(\mathbf{x}_N) & \phi_1(\mathbf{x}_N) & \dots & \phi_{M-1}(\mathbf{x}_N) \end{pmatrix}$$

ML Weights Linear Model:

$$\mathbf{w}_{ML} = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} = \Phi^\dagger \mathbf{t}$$



H3 Regularised Least Squares

Regularisation by introducing an **error term that penalises large weight values**:

$$\tilde{E}(\mathbf{w}) = E_D(\mathbf{w}) + \lambda E_W(\mathbf{w})$$

As before:

$$E_D(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N (t_n - \phi(\mathbf{x}_n)^T \mathbf{w})^2$$

H4 Ridge Regression

For **ridge regression**, weight penalty is $\lambda \|\mathbf{w}\|^2$, giving:

$$E_W(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} = \frac{1}{2} \|\mathbf{w}\|^2$$

where λ is the **regularisation coefficient**, controlling the relative importance of the two error terms. Total error function is now:

$$\begin{aligned} \tilde{E}(\mathbf{w}) &= \frac{1}{2} \sum_{n=1}^N (t_n - \phi(\mathbf{x}_n)^T \mathbf{w})^2 + \frac{\lambda}{2} \|\mathbf{w}\|^2 \\ &= \frac{1}{2} (\mathbf{t} - \Phi \mathbf{w})^T (\mathbf{t} - \Phi \mathbf{w}) + \frac{\lambda}{2} \mathbf{w}^T \mathbf{w} \end{aligned}$$

Differentiate, and set to zero:

$$\begin{aligned} \nabla_{\mathbf{w}} \tilde{E}(\mathbf{w}) &= -\Phi^T (\mathbf{t} - \Phi \mathbf{w}) + \lambda \mathbf{w} = 0 \\ -\Phi^T \mathbf{t} + \Phi^T \Phi \mathbf{w} + \lambda \mathbf{w} &= 0 \\ (\Phi^T \Phi + \lambda I) \mathbf{w} &= \Phi^T \mathbf{t} \\ \mathbf{w} &= (\Phi^T \Phi + \lambda I)^{-1} \Phi^T \mathbf{t} \end{aligned}$$

The error function for **ridge regression** (also known as **weight decay**) is quadratic which means it has a closed form solution too:

Regularised Weights Linear Model*:

$$\mathbf{w}^* = (\Phi^T \Phi + \lambda I)^{-1} \Phi^T \mathbf{t}$$

H4 Lasso

Other regularisation terms are also possible. For instance, **sum-of-absolute-values**:

$$E_W(\mathbf{w}) = \frac{1}{2} \sum_{j=1}^M |w_j|$$

- This approach is known as **lasso**
- If λ is sufficiently large, can lead to a **sparse model**: where most weight coefficients w_j are exactly zero.
- Sparse models can be more **robust** (resistant to over-fitting)

No general closed form solution for \mathbf{w}