

H1 Lecture 6: Probabilistic Classification - 24/02/20

H2 Generalised Linear Model

H3 Review: Discriminant Function

With linear regression predict real number with

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

A 2-class linear discriminant:

- **Evaluates** : $y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$
- **Assigns** \mathbf{x} to class C_1 if $y(\mathbf{x}) \geq 0$ and to C_0 otherwise
- Decision boundary defined by $y(\mathbf{x}) = 0$
- Decision boundaries are hyperplanes

To generalise this, we want a function $y(\mathbf{x})$ which...

- predicts class labels, $1, \dots, K$
- or posterior probabilities of class labels: $p(C_k|\mathbf{x})$
- e.g. for 2 -classes either $y(\mathbf{x}) \in \{0, 1\}$ or $y(\mathbf{x}) \in [0, 1]$

Can achieve this with **non-linear activation function** f and

$$y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x} + w_0)$$

- Decision boundaries are surfaces where $y(\mathbf{x}) = \text{constant}$, namely $(D - 1)$ - dimensional hyperplanes
- So decision surfaces are linear even though f is not

Generalised Linear Model

$$y(\mathbf{x}) = f(\mathbf{w}^T \mathbf{x} + w_0)$$

for non-linear activation function f

- Not linear in the parameters (unlike linear regression models)
- Implies more complex analytical and computational procedures
- Nevertheless, they are still relatively simple
- Can replace input, \mathbf{x} , with a fixed non-linear transformation to a vector of basis functions, $\Phi(\mathbf{x})$

H2 Generative Model


Consider the generative approach:

- class-conditional densities $p(\mathbf{x}|C_k)$
- class priors $p(C_k)$
- use Bayes Theorem to compute $p(C_k|\mathbf{x})$
- consider 2-classes:

$$\begin{aligned}
p(C_1|\mathbf{x}) &= \frac{p(\mathbf{x}|C_1)p(C_1)}{p(\mathbf{x}|C_1)p(C_1) + p(\mathbf{x}|C_0)p(C_0)} \\
&= \frac{p(\mathbf{x}|C_1)p(C_1)}{p(\mathbf{x}|C_1)p(C_1) + p(\mathbf{x}|C_0)p(C_0)} \times \frac{p(\mathbf{x}|C_1)p(C_1)}{p(\mathbf{x}|C_1)p(C_1)} \\
&= \frac{1}{1 + \frac{p(\mathbf{x}|C_0)p(C_0)}{p(\mathbf{x}|C_1)p(C_1)} \times \frac{p(\mathbf{x})}{p(\mathbf{x})}} \\
&= \frac{1}{1 + \frac{p(C_1|\mathbf{x})}{p(C_0|\mathbf{x})}} \\
&= \frac{1}{1 + \exp(-a(\mathbf{x}))}
\end{aligned}$$

where $a(\mathbf{x}) = \ln\left(\frac{p(C_1|\mathbf{x})}{p(C_0|\mathbf{x})}\right)$

H2 The Logistic Sigmoid



$$\sigma(a) = \frac{1}{1 + e^{-a}} \quad (6.2)$$

- Sigmoid shown in red
- Gaussian CDF shown in blue

- Sigmoid means S-shaped
- Maps from real line to interval $[0, 1]$
- Symmetric, $\sigma(-a) = 1 - \sigma(a)$
- Inverse called logit or log-odds-ratio:

$$a = \ln\left(\frac{\sigma}{1 - \sigma}\right) = \ln\left(\frac{p(C_1|\mathbf{x})}{p(C_0|\mathbf{x})}\right)$$

H2 2-Gaussian Classes with Identical Covariance


Assume each class distribution $p(\mathbf{x}|C_k)$ is Gaussian with the same covariance matrix, \mathbf{S} , then we can show

$$p(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

where

$$\begin{aligned}
\mathbf{w} &= \mathbf{S}^{-1}(\mu_1 - \mu_0) \\
w_0 &= -\frac{1}{2}\mu_1^T \mathbf{S}^{-1} \mu_1 + \frac{1}{2}\mu_0^T \mathbf{S}^{-1} \mu_0 + \ln \frac{p(C_1)}{p(C_0)}
\end{aligned}$$


- Posterior $p(C_1|\mathbf{x})$ is a generalised linear function
- Prior probabilities, $p(C_k)$ enter only through the bias term
- Minimum misclassification decision boundary is linear



Class-conditional densities for two classes:

$$p(\mathbf{x}|C_1) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_1, \mathbf{S}) \text{ in red}$$

$$p(\mathbf{x}|C_0) = \mathcal{N}(\mathbf{x}|\boldsymbol{\mu}_0, \mathbf{S}) \text{ in blue}$$



Corresponding posterior probability $p(C_1|\mathbf{x})$, given by a logistic sigmoid of a linear function of \mathbf{x} .

Assume a joint distribution for input \mathbf{x}_n , class C_k

$$p(\mathbf{x}_n, C_1) = p(C_1)p(\mathbf{x}_n|c_1) = \pi N(\mathbf{x}_n|\boldsymbol{\mu}_1, \mathbf{S})$$

$$p(\mathbf{x}_n, C_0) = p(C_0)p(\mathbf{x}_n|c_0) = (1 - \pi)N(\mathbf{x}_n|\boldsymbol{\mu}_0, \mathbf{S})$$

with prior $p(C_1) = \pi$, class means $\boldsymbol{\mu}_1$ & $\boldsymbol{\mu}_0$, and shared covariance \mathbf{S}

$$p(\mathbf{t}|\mathbf{X}, \pi, \boldsymbol{\mu}_1, \boldsymbol{\mu}_0, \mathbf{S}) = \prod_{n=1}^N q_1(\mathbf{x}_n)^{t_n} q_0(\mathbf{x}_n)^{(1-t_n)}$$

where

$$q_1(\mathbf{x}_n) = \pi N(\mathbf{x}_n|\boldsymbol{\mu}_1, \mathbf{S})$$

$$q_0(\mathbf{x}_n) = (1 - \pi)N(\mathbf{x}_n|\boldsymbol{\mu}_0, \mathbf{S})$$

Maximum Likelihood (ML) Solution for this model

Class-bias and means:

$$\pi^* = \frac{1}{N} \sum_n t_n = \frac{N_1}{N_1 + N_0}$$

$$\boldsymbol{\mu}_1^* = \frac{1}{N_1} \sum_n t_n \mathbf{x}_n$$

$$\boldsymbol{\mu}_0^* = \frac{1}{N_0} \sum_n (1 - t_n) \mathbf{x}_n$$

- N_k is number of points of class k
- Class bias, π , is the fraction of positive data-points
- Class-means are simply the means of each classes data-points

Covariance:

$$\mathbf{S}^* = \frac{N_1}{N} \mathbf{S}_1 + \frac{N_0}{N} \mathbf{S}_0$$

where $\mathbf{S}_1 = \frac{1}{N_1} \sum_n t_n (\mathbf{x}_n - \mu_1)(\mathbf{x}_n - \mu_1)^T$ and $\mathbf{S}_0 = \frac{1}{N_0} \sum_n (1 - t_n) (\mathbf{x}_n - \mu_0)(\mathbf{x}_n - \mu_0)^T$

- Covariance is the weighted sum of the class covariance
- Shared covariance assumed
- But it takes $\frac{D(D+1)}{2}$ computations for 2-class and $k \frac{D(D+1)}{2}$ for k -class, the time complexity is $O(D^2)$

H2 The Soft-Max

For K -classes:

$$\begin{aligned} p(C_k|\mathbf{x}) &= \frac{p(\mathbf{x}|C_k)p(C_k)}{\sum_{j=1}^K p(\mathbf{x}|C_j)p(C_j)} \\ &= \frac{\exp(a_k)}{\sum_{j=1}^K \exp(a_j)} \end{aligned}$$

A normalised exponential: a natural extension to the logistic sigmoid, and where $a_k = \ln p(\mathbf{x}|C_k)p(C_k)$

Normalised exponential is sometimes called **soft-max**, because:

- If $a_k \gg a_j$ for all $j \neq k$
- then $p(C_k|\mathbf{x}) \simeq 1$ and $p(C_j|\mathbf{x}) \simeq 0$

For K -classes all with the same covariance matrix:


$$a_k = \mathbf{w}_k^T \mathbf{x} + w_{k0}$$

where $\mathbf{w}_k = \mathbf{S}^{-1} \mu_k$ and $w_{k0} = -\frac{1}{2} \mu_k^T \mathbf{S}^{-1} \mu_k + \ln p(C_k)$

- terms a_k again linear in x
- minimum misclassification boundaries again linear

If each class has an **independent covariance matrix**

- a_k is quadratic in \mathbf{x}
- Gives rise to quadratic discriminant with quadratic decision boundaries



- Three classes: green, red & blue
- Left – class contours.
Green & red classes have the same covariance matrix
- Right – posterior class probabilities (given by colour density).
Minimum misclassification boundaries shown in white.

H2 A Discriminative Approach

Recall that we can directly learn posterior class probabilities $p(C_k|\mathbf{x})$, for classification:

- useful when complexities in $p(\mathbf{x}|C_k)$ do not (or only weakly) influence classification task
- or if no good distributional form for each $p(\mathbf{x}|C_k)$

We take inspiration from:

- 2-Classes with Gaussian densities and $\Sigma_1 = \Sigma_2$ leads to:

$$p(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + w_0)$$

- Logistic sigmoid of quadratic functions when $\Sigma_1 \neq \Sigma_0$

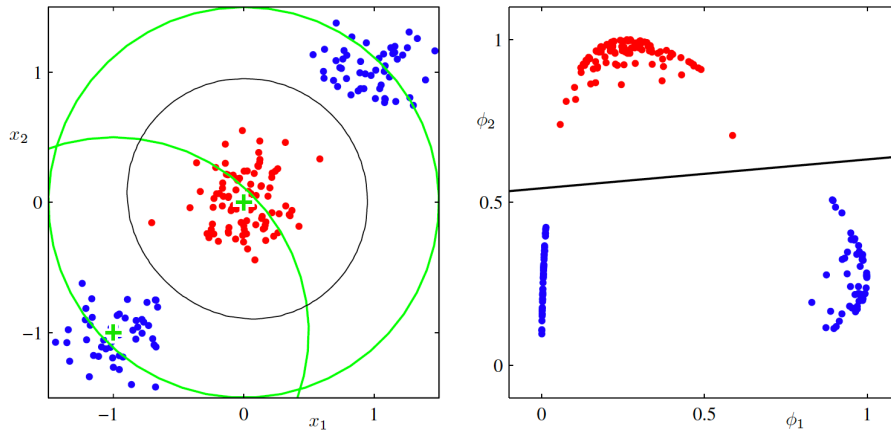
So why not instead, directly define a probabilistic discriminant using the logistic sigmoid?

- Typically this has fewer parameters to fit
- Can improve performance, if generative assumptions lead to poor approximations (e.g. classes not Gaussian)
- As with linear regression, we can use fixed basis functions:

$$\phi(\mathbf{x}) = (\phi_0(\mathbf{x}), \phi_1(\mathbf{x}), \dots, \phi_{M-1}(\mathbf{x}))$$

with $\phi_0(\mathbf{x}) = 1$

- No closed form solution, but iterative approach exist



- Left – classes not linearly separable & not Gaussian
- Left – 2 potential RBF centres and contours shown (green)
- Right – In feature space, classes are linearly separable

H2 Logistic Regression

Consider just 2-classes and fixed basis function ϕ :

- Posterior probability of C_1 written as:

$$p(C_1|\mathbf{x}) = \sigma(\mathbf{w}^T \phi(\mathbf{x}))$$

(No separate bias term, as $\phi(\mathbf{x}) = 1$)

- Using logistic sigmoid:

$$\sigma(a) = \frac{1}{1 + e^{-a}}$$

- For M basis functions, we have M -parameters (elements of \mathbf{w})
- Comparison: to fit our generative model our parameters would comprise: $2M$ for the means, $\frac{M(M+1)}{2}$ for shared covariance and 1 for class bias

H2 Discriminative Likelihood and Related Error

For data-points (ϕ_n, t_n) where $t_n \in \{0, 1\}$ and $\phi_n = \phi(\mathbf{x}_n)$ for $n = 1, \dots, N$ our discriminative model has likelihood:

$$p(\mathbf{t}|\mathbf{w}) = \prod_{n=1}^N y_n^{t_n} (1 - y_n)^{1-t_n}$$

for $\mathbf{t} = (t_1, \dots, t_N)^T$ and $y_n = p(C_1|\phi_n) = \sigma(\mathbf{w}^T \phi_n)$

As with regression, take negative log likelihood as error function:

The Cross-Entropy Error Function

$$\begin{aligned}
 E(\mathbf{w}) &= -\ln p(\mathbf{t}|\mathbf{w}) \\
 &= -\sum_{n=1}^N (t_n \ln y_n + (1 - t_n) \ln(1 - y_n))
 \end{aligned}$$

Take gradient to minimise the error:

$$\begin{aligned}
 \nabla_{\mathbf{w}} E(\mathbf{w}) &= \sum_{n=1}^N (\sigma(\mathbf{w}^T \Phi_n) - t_n) \Phi_n \\
 &= \Phi(\mathbf{y} - \mathbf{t})
 \end{aligned}$$

with design-matrix, Φ , targets $\mathbf{t} = (t_1, \dots, t_N)^T$, predictions $y_n = \sigma(\mathbf{w}^T \phi_n)$ and prediction vector $\mathbf{y} = (y_1, \dots, y_N)^T$

Unfortunately, the non-linear form means we cannot simply set to zero and rearrange

H2 Gradient Ascent Methods

We want to find the maximum of function $f: \mathbb{R}^D \rightarrow \mathbb{R}$, and we can calculate the gradient $\nabla f = \nabla_{\mathbf{z}} f$ for any point \mathbf{z}


Recipe:

- Choose an initial estimate $\mathbf{z}^{(0)}$ (possibly randomly)
- Repeatedly update estimate with:

$$\mathbf{z}^{(\tau+1)} = \mathbf{z}^\tau + \eta \nabla f$$


where $\eta > 0$ is a small step size and ∇f is evaluated at $\mathbf{z}^{(\tau)}$

- Stop when magnitude of ∇f falls below a threshold
- Eventually reaches a local maximum, if η small enough



H3 Problem

- Contours of f in orange
- Gradient ascent only finds local-maxima
- Path may not be direct (zig-zags)
- Must choose step-size, η



H2 Newton-Raphson Method

Consider a univariate function $f : \mathbb{R} \rightarrow \mathbb{R}$ which is twice differentiable. Define f_τ , the second order Taylor expansion of f around point $x^{(\tau)}$:

$$\begin{aligned} f(x) &\approx f_\tau(x) = f_\tau(x^{(\tau)} + u_\tau) \\ &= f(x^{(\tau)}) + f'(x^{(\tau)})u_\tau + \frac{1}{2}f''(x^{(\tau)})u_\tau^2 \end{aligned}$$

where $u_\tau = x - x^{(\tau)}$

u_τ Maximises/minimises this expression when:

$$\begin{aligned} \frac{d}{du_\tau}(f_\tau(x^{(\tau)} + u_\tau)) &= f'(x^{(\tau)}) + f''(x^{(\tau)})u_\tau = 0 \\ \implies u_\tau &= -\frac{f'(x^{(\tau)})}{f''(x^{(\tau)})} \end{aligned}$$

Message: can use second order derivatives to climb functions more quickly

We seek the maximum of function $f : \mathbb{R}^D \rightarrow \mathbb{R}$, and can calculate gradient $\nabla_z f$ and Hessian $\mathbf{H}(\mathbf{z}) = \nabla \nabla f = \nabla^2 f$ for any \mathbf{z} .

The Hessian is the matrix of second derivatives evaluated at \mathbf{z} :

$$[\mathbf{H}(\mathbf{z})]_{ij} = \frac{\partial^2 f}{\partial z_i \partial z_j} \Big|_{\mathbf{z}}$$

$$\mathbf{H}(\mathbf{z}) = \begin{pmatrix} \frac{\partial^2 f}{\partial z_1^2} & \frac{\partial^2 f}{\partial z_1 \partial z_2} & \cdots & \frac{\partial^2 f}{\partial z_1 \partial z_n} \\ \frac{\partial^2 f}{\partial z_2 \partial z_1} & \frac{\partial^2 f}{\partial z_2 \partial z_2} & \cdots & \frac{\partial^2 f}{\partial z_2 \partial z_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial z_n \partial z_1} & \frac{\partial^2 f}{\partial z_n \partial z_2} & \cdots & \frac{\partial^2 f}{\partial z_n^2} \end{pmatrix}$$

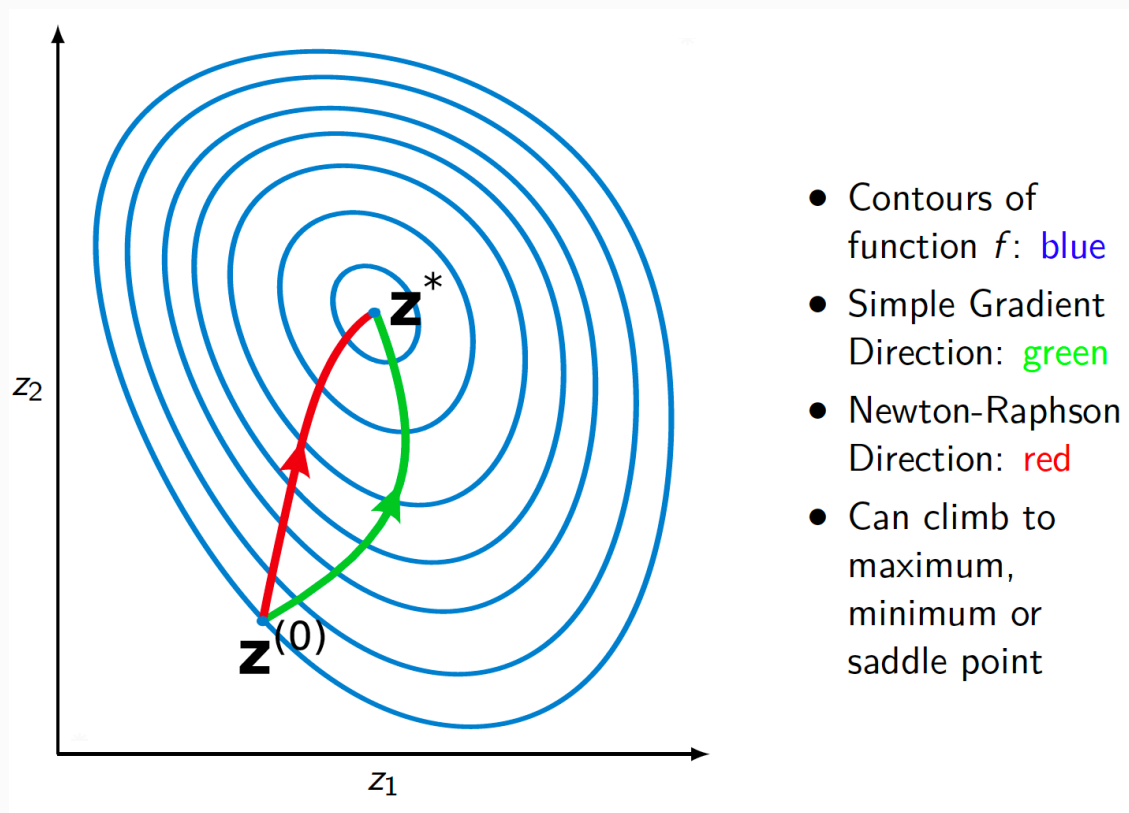
Recipe:

- Choose an initial estimate $\mathbf{z}^{(0)}$ (possibly randomly)
- Repeatedly update estimate with:

$$\mathbf{z}^{(\tau+1)} = \mathbf{z}^{(\tau)} - \mathbf{H}^{-1} \nabla f$$

where ∇f and \mathbf{H} are evaluated at $\mathbf{z}^{(\tau)}$

- Stop when magnitude of $\mathbf{H}^{-1} \nabla f$ falls below a threshold



H2 Linear Regression with Newton-Raphson Method

Apply the **Newton-Raphson Method** to Linear Regression model:

$$\begin{aligned} E(\mathbf{w}) &= \frac{1}{2} (\mathbf{t} - \Phi \mathbf{w})^T (\mathbf{t} - \Phi \mathbf{w}) \\ \nabla E(\mathbf{w}) &= \Phi^T \Phi \mathbf{w} - \Phi^T \mathbf{t} \\ \mathbf{H}(\mathbf{w}) &= \nabla^2 E(\mathbf{w}) = \Phi^T \Phi \end{aligned}$$

Iterative function:

$$\begin{aligned} \mathbf{w}^{(new)} &= \mathbf{w}^{(old)} - (\Phi^T \Phi)^{-1} (\Phi^T \Phi \mathbf{w}^{(old)} - \Phi^T \mathbf{t}) \\ &= (\Phi^T \Phi)^{-1} \Phi^T \mathbf{t} \\ &= \mathbf{w}_{ML} \end{aligned}$$

We get the maximum in a **single step** for any $\mathbf{w}^{(old)}$

H2 Logistic Regression with Newton-Raphson Method

For Logistic Regression, we seek to minimise **cross-entropy error**:

$$E(\mathbf{w}) = - \sum_{n=1}^N (t_n \ln y_n + (1 - t_n) \ln(1 - y_n))$$
$$\nabla E(\mathbf{w}) = \Phi(\mathbf{y} - \mathbf{t})$$

with design-matrix, Φ , targets $\mathbf{t} = (t_1, \dots, t_N)^T$, predictions $y_n = \sigma(\mathbf{w}^T \phi_n)$ and prediction vector $\mathbf{y} = (y_1, \dots, y_N)^T$

Hessian Matrix:

$$\mathbf{H} = \nabla \nabla E(\mathbf{w}) = \sum_{n=1}^N y_n(1 - y_n) \Phi_n^T \Phi_n = \Phi^T \mathbf{R} \Phi$$

Where $\mathbf{R} \in \mathbb{R}^{N \times N}$ is a diagonal matrix with non-zero elements:

$$[\mathbf{R}]_{nn} = y_n(1 - y_n)$$

H3 Algorithm

Iteratively reweighted least squares (IRLS) algorithm [Rub83]:

```
1: procedure IRLS( $\Phi$ ,  $\mathbf{t}$ ,  $\mathbf{w}^{(0)}$ ,  $\theta$ )
2:   #  $\Phi$  is design matrix,  $\mathbf{t}$  is target vector
3:   #  $\mathbf{w}^{(0)}$  is initial weight vector,  $\theta > 0$  is threshold
4:   for  $\tau = 1, 2, \dots$  do
5:     for  $n = 1, \dots, N$  do
6:        $y_n = \sigma(\mathbf{w}^T \phi_n)$ 
7:        $r_{nn} = y_n(1 - y_n)$ 
8:      $\mathbf{y} = (y_1, \dots, y_N)^T$ 
9:      $\mathbf{R} = \text{diag}((r_{11}, \dots, r_{nn})^T)$ 
10:     $\mathbf{w}^{(\tau)} = \mathbf{w}^{(\tau-1)} - (\Phi^T \mathbf{R} \Phi)^{-1} \Phi^T (\mathbf{y} - \mathbf{t})$ 
11:    if  $\|\mathbf{w}^{(\tau)} - \mathbf{w}^{(\tau-1)}\| < \theta$  then
12:      return  $\mathbf{w}^{(\tau)}$ 
```

where ∇E in red, \mathbf{H} in blue, $\text{diag}(\cdot)$ constructs a diagonal matrix from a vector, and $\|\mathbf{w}^{(\tau)} - \mathbf{w}^{(\tau-1)}\|$ is Euclidean distance from $\mathbf{w}^{(\tau)}$ to $\mathbf{w}^{(\tau-1)}$.

Important: \mathbf{y} and \mathbf{R} are re-evaluated each iteration